


Robustness gap under translation attack



Bold lines: aggregate TPR across 3 LLMs x 4 domains x 100 prompts ($z_{\alpha} = 2.326$). Thin transparent lines: per-(model, domain) cells.